

Your repo's contribution rules, compiled into an agentic gate.

The deliverable is Bob configuration — a contract Bob compiled from the repository's own documents, three custom modes, twelve rules files — and the deterministic checkers those modes drive.

Target repository **ScrollPrize/villa** · evaluated on ten real pull requests · github.com/CVasilopoulos/reviewready

Four working fixes. **Sixty-one days open.** Nobody read them.

PR	WHAT IT FIXED	OPEN FOR	OUTCOME
#1391	vc_obj2tifxyz TIFXYZ scale semantics	17 days	closed unreviewed
#1412	test runner: abort the test slot	14 days	closed unreviewed
#1630	fail instead of writing sentinels	16 days	closed unreviewed
#1724	tifxyz_label_transfer: close memory maps	14 days	closed unreviewed

“reviewing PRs takes a significant amount of time away from the goal of unrolling the scrolls”

ScrollPrize/villa · CONTRIBUTING.md, AI Guidelines

Two of the four were later fixed again, by other people, from scratch. The rules that would have saved all of this are already written down — in prose, where nothing enforces them.

The repo already runs two gates. One counts days. One counts files.

PR-TIME-LIMITS.YML

```
inactive_cutoff=$(date -u -d '14 days ago' +%s)
age_cutoff=$(date -u -d '28 days ago' +%s)
...
gh pr close "$number"
```

Fires on four of our ten. Every time by **closing a working fix nobody read**, telling its author nothing.

LARGE-PR-REVIEW-GATE.YML

```
LARGE_PR_FILE_THRESHOLD: "20"
// block large PRs unless approved
```

Fires on none of them. No pull request in the corpus touches twenty files.

The three questions that actually decide a pull request's fate — has someone already fixed this? does the body satisfy the written rules? is the proof really a proof? — are reading problems, not counting problems. That is the gap an agent fills and a workflow cannot.

Prose in, **executable contract** out.

CONTRIBUTING.MD — WHAT A HUMAN WROTE

*"Any bugfix PR must be accompanied by a screenshot of the error ...
and the script/tool running without error afterward"*

line 19, one of five bullets

CONTRACT.YML — WHAT BOB COMPILED


```
- id: bugfix-error-screenshot
  source: CONTRIBUTING.md:19
  check: (a) a screenshot or terminal block
         showing the error, and (b) one showing
         the tool running without error after
  severity: blocker
```

THE RULE NOBODY WROTE DOWN — BOB FOUND THE HOLE

villa has six rules for LLM-assisted PRs and **no rule that you must disclose using one**. Without disclosure all six are unenforceable. Bob emitted **llm-disclosure** as source: `inferred`, carrying a `rationale` field instead of a quote — so a maintainer can see it was derived, and reject it.

27 rules. 25 quote a source line. 2 are marked `inferred`. Nothing is hard-coded in a checker. Before compiling, Bob fanned out **four subagents in parallel** to survey villa's subprojects — 21 minutes, 3.69 Bobcoins, every call logged in `bob_sessions/01`.

Three custom modes. Read, execute, todo — nothing else.



No mode may write, edit or merge. Each loads four rules files before touching anything — role, protocol, output schema, limits. **Five modes were planned; three exist** — the Bobcoin allocation ran out after the third.

#1797 — the gate read the rulebook. The maintainer read the design.

WHAT REVIEWREADY SAID

REVISE

`llm-concise-description` — description too long.

`llm-human-review` — no statement of having read the diff.

WHY IT WAS ACTUALLY CLOSED

"Remote rendering should either take a remote path ... or a local volume -v. I see no good reason to add this mixture?"

maintainer, 2026-09-22 · PR closed

*Nothing in CONTRIBUTING.md says "your flags must compose into a coherent interface", so no rule in the contract can. This is the fourth checker — **interface shape** — and it does not exist: the budget ran out at three. We report it as a miss because it marks exactly where a compiled contract stops and a human reviewer begins.*

MEASURED, ON PULL REQUESTS WE DID NOT WRITE

Five of five caught. Zero good ones blocked.

5 / 5

HOLD-OUT RECALL

0 / 2

HOLD-OUT FALSE BLOCKS

6 / 7

WHOLE CORPUS RECALL

1

FALSE BLOCK, WHOLE CORPUS

PR	GATE'S TOP REASON	WHAT REALLY HAPPENED	AGREE	INCUMBENT CI
#1391	real-scroll-data-origin	auto-closed at 17d, 0/6 template fields	yes	closed it
#1683	prior-art/closed-pr-same-issue	duplicate of #1391, closed by author	yes	nothing
#1794	prior-art/closed-pr-same-issue	superseded eight minutes later	yes	nothing
#1724	verification-checkbox (major)	fully compliant — auto-closed at 15d unread	yes	closed it
#1797	llm-concise-description	closed on interface design	miss	nothing
#1798	evidence/no-before-commit	clean PR — our false positive	no	nothing

Neither incumbent gate catches either duplicate. Both duplicates were written because the earlier fix had been closed and was invisible. #1798 fails because its commit hash sits 800 characters from the nearest before/after label; widening the window further would be fitting the checker to this corpus, so it is left failing. **Tuned on our own PRs, evaluated on other people's. Every row is a public PR you can open.** Verdicts here come from running the three checker scripts over the cached corpus; the Bob modes drive those same scripts.

A folder you drop into any repository. Then one prompt re-derives it.

THE GATE PACK

<code>.bob/custom_modes.yaml</code>	three modes, minimal tool grants
<code>.bob/rules-<slug>/</code>	twelve rules files
<code>contract.yaml</code>	27 rules, each with its source line
<code>scripts/check_*.py</code>	the script each mode drives
<code>schemas/finding.json</code>	one finding format
<code>reports/</code>	ten PRs, scored, with SUMMARY.md

- Nothing is villa-specific except `contract.yaml` — and Bob wrote that by reading villa's own documents.
- Point the same prompt at another repo's `CONTRIBUTING.md` and PR template, get that repo's contract.
- Runs before you open the pull request, not after fourteen silent days. The three checkers take under a second on a cached PR; prior-art costs a handful of gh calls.